Discriminating Sound Textures

Jayson Lynch* Eric Mannes†

December 9, 2015

Abstract

We show sound textures are cool stuff.

Keywords: Machine Learning, Sound Textures, Audition

*MIT Computer Science and Artificial Intelligence Laboratory, 32 Vassar Street, Cambridge, MA 02139, USA, {jaysonl}@mit.edu.

†MIT Computer Science and Artificial Intelligence Laboratory, 32 Vassar Street, Cambridge, MA 02139, USA, {mannes}@mit.edu.

Contents

1 Introduction

Stuff [Cite McDermott's papers [MS11, MSS13] and why we care about these features/classification] xxx

[Cite some other audio/sound texture classification work.] xxx

2 Definitions

[Give formal and informal definitions of the features we used] xxx

3 Methods

[Describe how we set up our feature extraction. How we set up our learning. What our dataset is. Any other important process things.] xxx

4 Results

[What coorilated and what didn't? What feature (ensemble) lead to a good classifier?] xxx

5 Conclusion

[Speculate if these could be useful features in audio classification. Speculate about the implication to human audition. Give future directions.] xxx

Acknowledgments

We thank Prof Josh McDermott for his support in answering our questions about their research and providing us with their dataset. We also thank the course staff of 6.867 for their instruction and support this semester. In particular, we appreciate Marzyeh Ghassemi advice and guidance in shaping our project plan.

References

[MS11] Josh H. McDermott and Eero P. Simoncelli. Sound texture perception via statistics of the auditory periphery: Evidence from sound synthesis. *Neuron*, 71(5):926 – 940, 2011.

[MSS13] Josh H McDermott, Michael Schemitsch, and Eero P Simoncelli. Summary statistics in auditory perception. *Nature neuroscience*, 16(4):493–498, 2013.